

Optimizing Data File Layout

Learning Objectives

- ▶ The concept of data file layout
- ▶ Exploring optimization techniques.

Data File Layout

- ▶ The organization and storage structure of the underlying data files that make up a Delta table.
- ▶ Optimizing layout helps leveraging data-skipping algorithms
- ▶ Optimization techniques:
 - ▶ Partitioning
 - ▶ Z-Order Indexing
 - ▶ Liquid Clustering

Partitioning

- Organizing a table by grouping rows that share the same values for predefined partitioning columns

partitioning column
↓

id	name	month	year
1	Adam	9	2015
2	Sarah	11	2015
3	John	5	2000
4	Ali	9	2000
5	Emma	10	2000
6	Eric	1	2023
7	Sophia	3	2023

Partition 1
year=2015

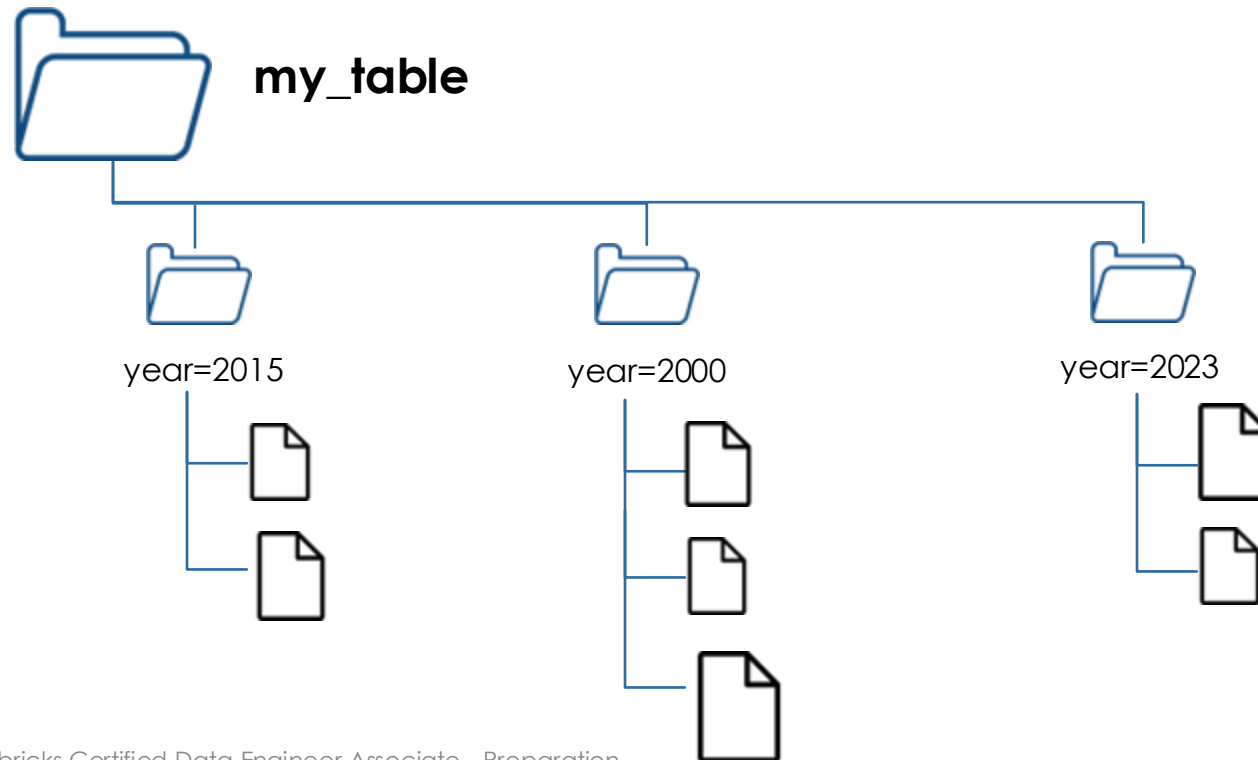
Partition 2
year=2020

Partition 3
year=2023

Partitioning Delta Lake Tables

CREATE TABLE my_table (id INT, name STRING, year INT, month INT)

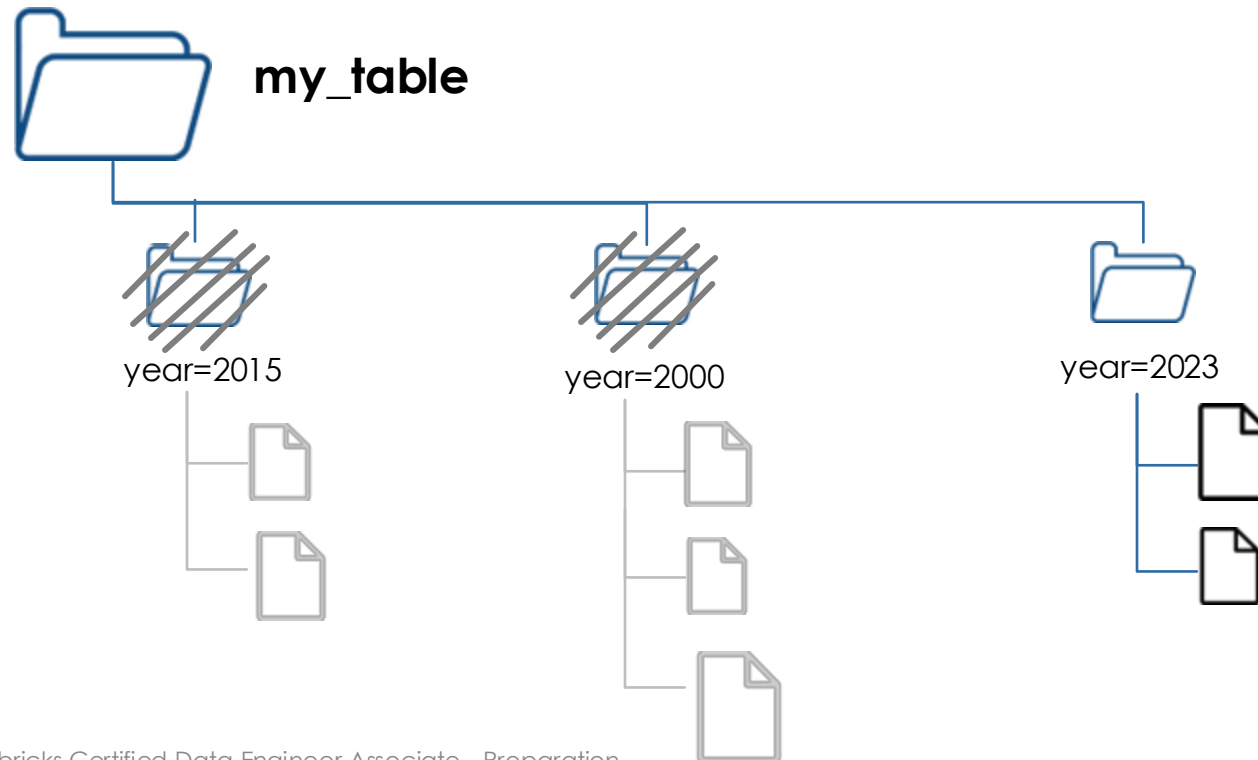
PARTITIONED BY (year)



Partition Skipping

```
SELECT * FROM my_table
```

Where year = 2023




Partitioning Limitations

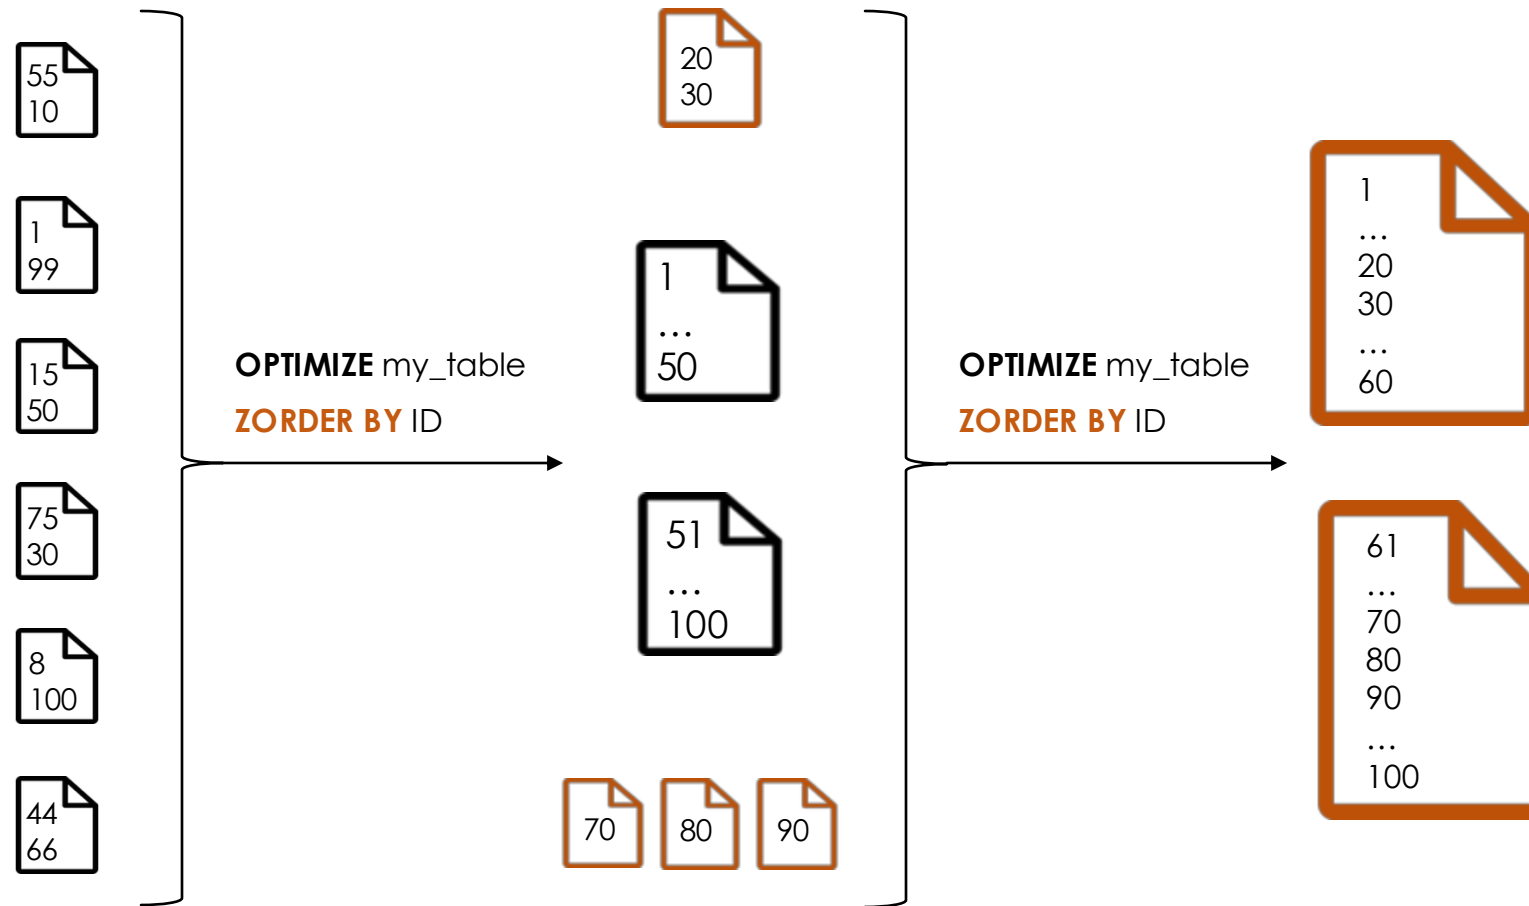
- ▶ Prevents file compaction across partition boundaries
 - ▶ Results in a small files problem
- ▶ Inefficient for high-cardinality columns
 - ▶ Results in a small files problem
- ▶ Static: Re-partitioning requires a full table rewrite

Z-Order Indexing

- ▶ Group similar data into optimized files without creating directories
- ▶ Leverage data-skipping algorithms
- ▶ Effective for High-cardinality columns



Z-Ordering: Not Incremental



Liquid Clustering

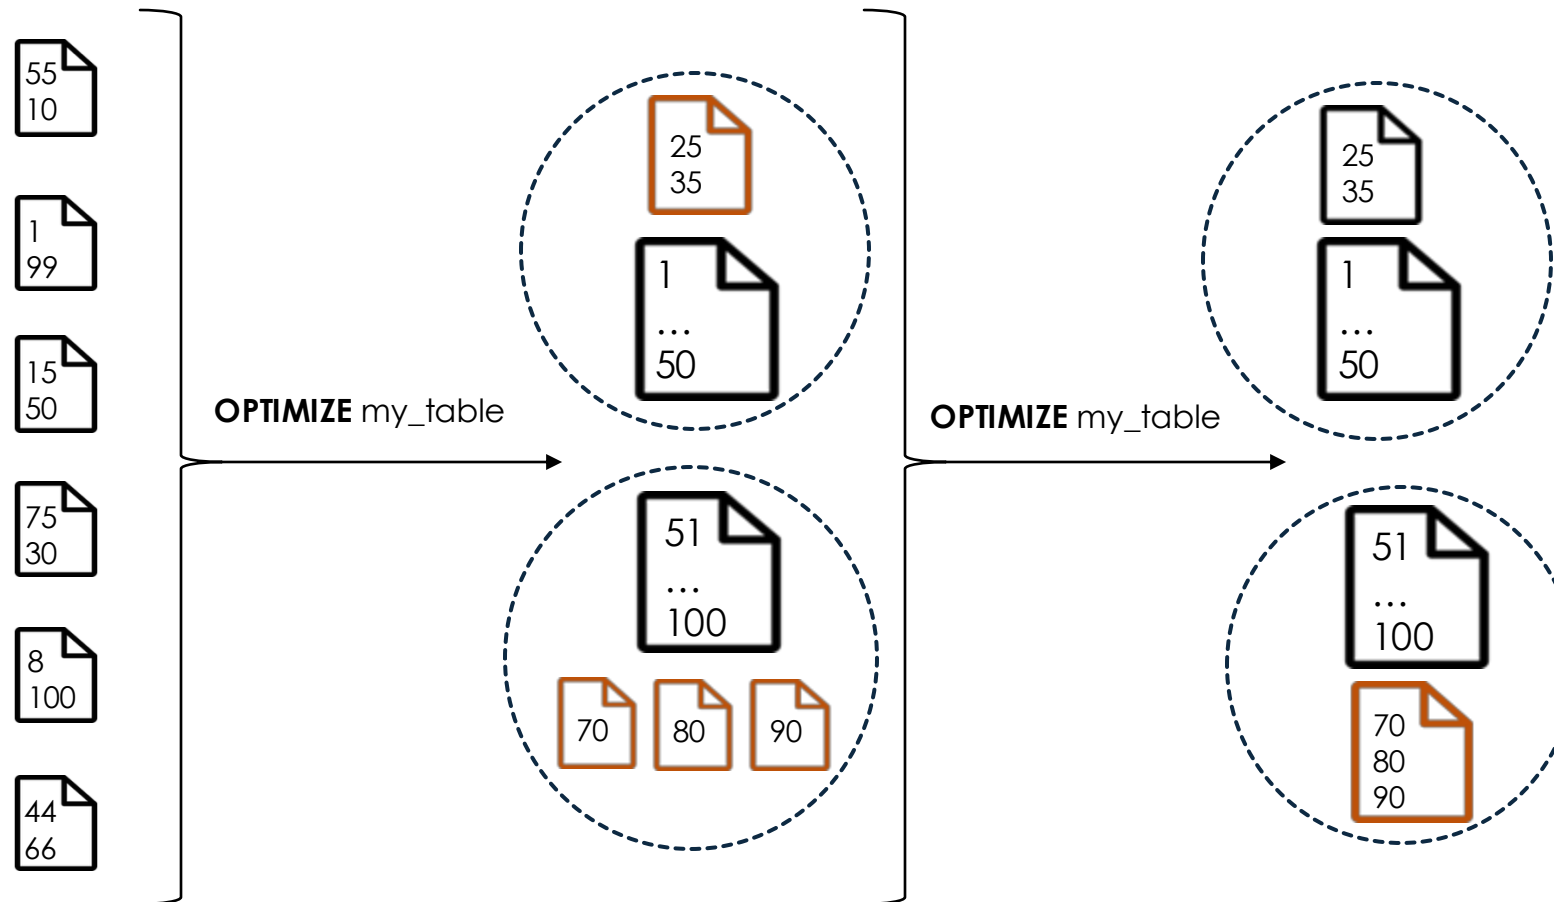
- ▶ Improved version of Z-order indexing with more flexibility and better performance.
- ▶ Table-level definition
 - ▶ New tables:

```
CREATE TABLE table1(col1, INT, col2 STRING, col3 DATE)  
CLUSTER BY (col1, col3)
```
 - ▶ Existing tables:

```
ALTER TABLE table2  
CLUSTER BY (<clustering_columns>)
```
- ▶ Clustering is not compatible with partitioning or ZORDER

Incremental Clustering

```
ALTER TABLE my_table  
CLUSTER BY ID
```



Choosing Clustering Keys

- ▶ Flexible to redefine clustering keys without rewriting existing data
- ▶ Choose clustering keys based on your query pattern

Automatic Liquid Clustering

- ▶ Databricks automatically chooses clustering keys by analyzing the table historical query workload.
- ▶ Requires Predictive Optimization on Unity Catalog managed tables

- ▶ Syntax

- ▶ New tables:

```
CREATE TABLE table1(col1, INT, col2 STRING, col3 DATE)  
CLUSTER BY AUTO
```

- ▶ Existing tables:

```
ALTER TABLE table2  
CLUSTER BY AUTO
```